

# Research projects

Li Wenliang (Kevin)

Gatsby Computational Neuroscience Unit, University College London  
[kevin-w-li.github.io](https://github.com/kevin-w-li)

April 4, 2018

# Topics

- ▶ General interest
  - ▶ How the brain works in perception and cognition.
  - ▶ Use machine learning as a tool to study neuroscience.
  - ▶ Inspire machine learning algorithms from studying neuroscience.
- ▶ Theoretical neuroscience
  - ▶ Perceptual learning: Wenliang and Seitz, in press
  - ▶ Neural representation of uncertainty:
    - ▶ Wenliang and Sahani, 2018
- ▶ Machine learning
  - ▶ Kernel methods (Arthur Gretton)
    - ▶ density estimation
    - ▶ goodness-of-fit test, on-going
  - ▶ Approximate inference (Maneesh Sahani), on-going

# Perceptual learning

## **Simple (Bayesian) networks**

- ▶ Data: Doshier and Lu (1998)
- ▶ Method: learning as inference
- ▶ Unexpected insight into where learning occurs?
  - ▶ How do the brain learns to classify patterns of stimulus representation (high areas)
  - ▶ How do the representations themselves adapt to the task? (low areas)

## **Deep neural networks**

- ▶ Factors that modulate transfer
- ▶ Distribution of learning over a deeper hierarchy
  - ▶ How does this distribution change w.r.t. task?
  - ▶ How much does the layers contribute to performance?
- ▶ Similarities to physiological studies (e.g. Schoups et al. 2001)

## **Unsupervised perceptual learning (on going)**

## Bayesian network: ignoring subcortical

**Stimulus:** Gabors at  $+s/2$  and  $-s/2$  generate neural representations  $\mathbf{x}_{-1}$  and  $\mathbf{x}_{+1}$

$$y \in \{-1, +1\} \sim \text{Bernoulli}(0.5)$$

$$\mathbf{x}|y \sim \mathcal{N}(c\bar{\mathbf{x}}_y, \mathbf{\Sigma}(c, \bar{\mathbf{x}}_y))$$

$$\bar{\mathbf{x}}_{y,a} = \exp \left[ -\frac{(a + ys/2)^2}{2} \right]$$

**Discrimination:** noisy linear classification using probabilistic weights

$$\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \sigma_w^2 \mathbf{I})$$

$$b = \phi(\mathbf{w} \cdot \mathbf{x})$$



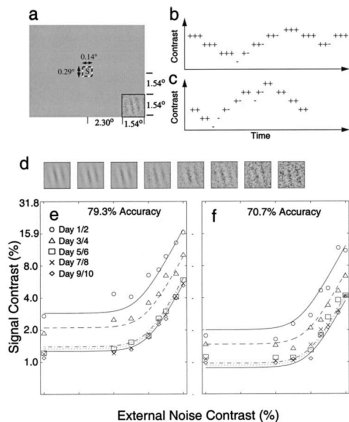
$$\hat{y} = \text{Bernoulli}(b)$$

**Learning:** update  $\mathbf{w}$  given training data  $\{\hat{y}_i, \mathbf{x}_i\}$

$$p(\mathbf{w} | \{\hat{y}_i, \mathbf{x}_i\}_{i=1}^k) \propto p(\mathbf{w}) \left[ p(\{\hat{y}_i = y_i\}_{i=1}^k | \mathbf{w}) \right]$$


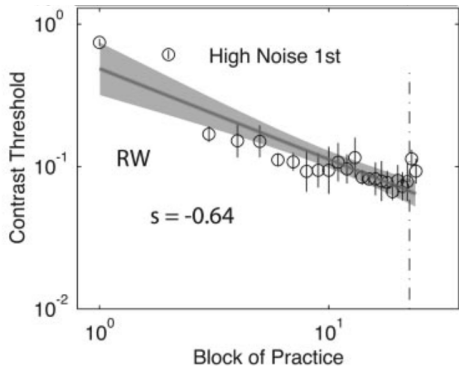
# Bayesian network: learning in high areas

Doshier and Lu 1998



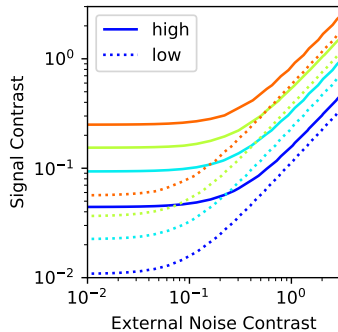

# Bayesian network: results

Doshier and Lu 2005



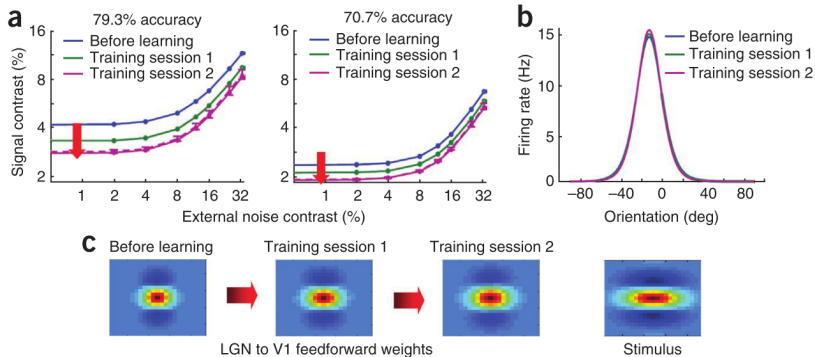
# Bayesian network: asymmetric transfer

Doshier and Lu 2005



# However, something is missing

- ▶ This model so far is able to qualitatively reproduce the results modelled by Doshier and Lu (2010)
- ▶ However, Bejjanki et al. (2011) offered another perspective: the TVC curves can be achieved by changes in subcortical connections



- ▶ Which is correct?



# Taking into account sub-cortical pathways


- ▶ Now we add noise directly to the images.

$$\begin{aligned} \mathbf{I}_y &\sim N(c\bar{\mathbf{I}}_y, \sigma_n^2 \mathbb{I}) \\ x_a|y &\sim \text{Poisson}(\mathbf{v}_a^T \mathbf{I}_y) \\ b &= \phi(\mathbf{w} \cdot \mathbf{x}) \end{aligned}$$

- ▶ The representations  $\mathbf{x}$  are correlated, and changing only  $\mathbf{w}$  no longer reproduce the same effect on TVC.
- ▶ However, by re-aligning  $\mathbf{v}$  towards  $\mathbf{I}$ , we can reproduce the effect shown by Bejjanki et al. (2011).
- ▶ Note that, in Doshier and Lu (2010), independent noise are added at many intermediate stages, effectively reducing the noise correlations of the activities used for classification.
- ▶ **Hypothesis:**
  - ▶ if neuronal noise makes the representations independent, then only higher level changes can explain data;
  - ▶ if not, then the representations have to change.

# Deep neural network model

- ▶ Recent similarities found in the literature (Yamins, Kriegeskorte, etc) motivates the use of DNN to investigate VPL.
- ▶ We set up a deep neural network to simulate learning of Gabor orientation discrimination under 2AFC.
- ▶ The weights are initialized from the first five layers of pre-trained AlexNet

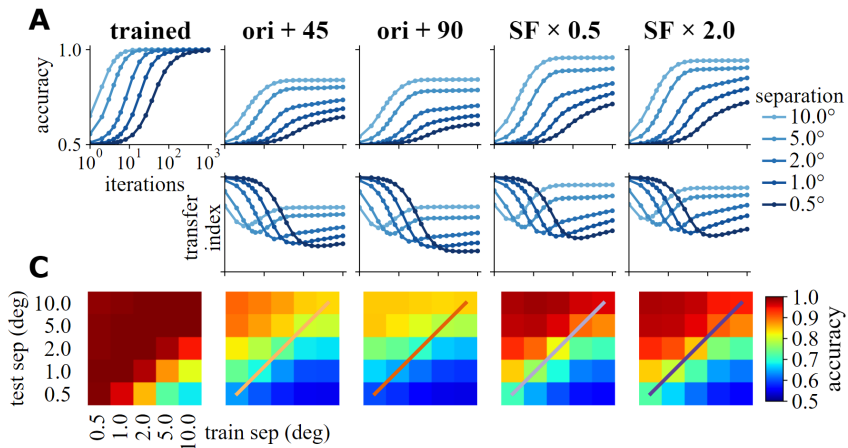


- ▶ **Questions: can this model reproduce findings of VPL?**

# DNN model: Behavioral level

Behavioral performance measured as percentage correct



- Specificity prevails for more precise discrimination (Ahissar and Hochstein 1997)
- Test precision has a major effect while training precision may help transfer for coarser discrimination



## DNN model: System level

Distribution of weight change moves towards lower layers for

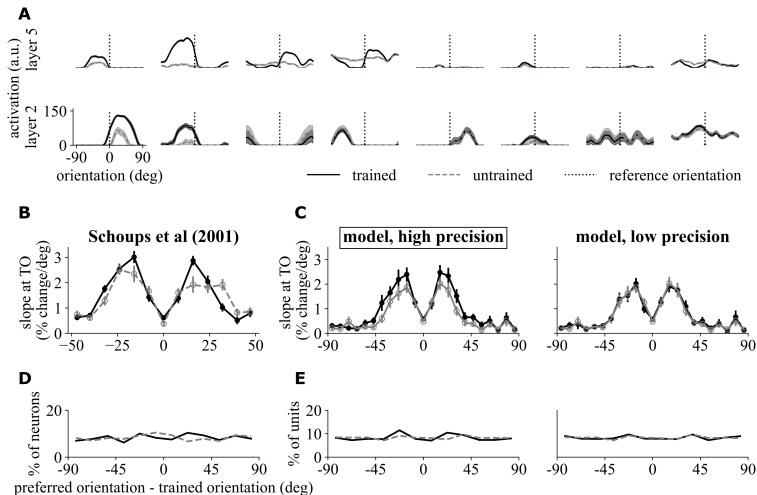
- ▶ more precise discriminations
- ▶ lower-level task (orientation vs gender)



# DNN model: Neuronal level

Explore similarities between visual neurons in monkeys to the units in the DNN

- Qualitatively match without direct fitting to data




- Also reproduces double training (but with convolution)

# Unsupervised perceptual learning


- ▶ Experiments show that perceptual learning can happen without supervision signal
  - ▶ By exposure along with a irrelevant task and stimulus (Watanabe et al. 2001)
  - ▶ By pairing with reward (Seitz et al. 2009)
  - ▶ By mental imagery (Tartaglia 2009)
  - ▶ By passive exposure in absence of attention (Amitay 2006)
- ▶ Conjectures and Approaches
  - ▶ Use unsupervised learning methods to obtain features/representations (ICA)
  - ▶ Explanation by long-term adaptation (Harris et al. 2012)

# Neural representation of uncertainty


- Humans and animals are able to process uncertain information as if they are able to manipulate the underlying probability distributions




- How might the brain achieve this? Three problems need to be solved



Representation



Computation




Learning

# Distributed distributional representation (DDC)

We propose the DDC scheme for representing and computing with distributions

- ▶ The mean rate of a neuron encode the expectation of some nonlinear functions  $\psi(\cdot)$  under the probability distribution encoded.



- ▶ Expectations define the underlying distribution by maximum entropy,
  - ▶ e.g. for Gaussians, only need:  $[r_1, r_2] = [\mathbb{E}_{p(y)}[y], \mathbb{E}_{p(y)}[y^2]]$
  - ▶ In general, these expectations define the **mean parameters** of an exponential family distribution

$$\mu_y := [r_1, r_2, \dots, r_k] \xrightarrow{\text{max. ent.}} p(y) \propto \exp \left( \sum_i \theta_i \psi_i(y) \right)$$



# Computation and learning with DDC

Usually, the quantity of interest is the **expectation** of some other function  $g(y)$

- ▶ In prediction/regression, the quantity that minimises the squared loss is the **expectation** under the predictive distribution.
- ▶ In reinforcement learning, the value function  $Q(s, a)$  is an **expectation** over the world

Using DDC, approximating an expectation is straightforward

$$g(y) \approx \sum_i \alpha_i \psi_i(y) \Rightarrow \mathbb{E}_{p(y)} [g(y)] \approx \sum_i \alpha_i \mathbb{E}_{p(y)} [\psi_i(y)] = \sum_i \alpha_i r_i$$

Bayes rule in densities corresponds to linear regression in DDC

$$p(y|x) \propto p(y)p(x|y) \Leftrightarrow \mathbb{E}_{p(y|x)}[\psi(y)] = \mathcal{W}\psi(x)$$


where  $\mathcal{W} = \mathbb{E}[\psi(y) \otimes \phi(x)] \mathbb{E}[\phi(x) \otimes \phi(x)]^{-1}$  estimated from training data

# Automatic appearance in neural network

A tracking task: estimate  $y_t$  given all available data  $x_{1:t}$

$$y_t \sim \text{Normal}(f(y_{t-1}), \sigma^2)$$

$$x_t \sim \text{Poisson}(g(y))$$




To minimise the mean squared error, the solution is  $\mathbb{E}_{p(y_t|x_{1:t})} [y_t]$


Inference requires belief updating (c.f. kernel Bayes rule):

$$p(y|x_{1:t}) \propto p(y|x_{1:t-1})p(y_t|y_{t-1})p(x_t|y_t) \Leftrightarrow \mu_y = \sum_{jk} W_{ijk} \mu_{y-1,j} \phi_k(x)$$

This defines a Bilinear RNN. So far there has been no constructions of  $\psi(y)$




Generative



Bilinear

# Automatic appearance in neural network

The DDC neural network should perform well on the task





The neurons adopt DDC representation

- activation equals posterior expectation of some nonlinear function  $\psi(y_t)$

$$r_t^{(i)} \stackrel{?}{=} \mathbb{E}_{p(y_t | x_{1:t})} [\psi_i(y_t)]$$

- $\psi(y_t)$  found by function approximation (RBF)



# Machine learning

- ▶ Kernel methods for
  - ▶ density estimation: score matching with a neural network
  - ▶ goodness of fit: benchmarking implicit models

# Density estimation with exponential family distributions

Fit a probability distribution for data  $x \sim p(x)$

- ▶ log density function: infinite dimensional exponential family

$$q(x) = \frac{1}{Z} \exp[f(x)]$$

- ▶ using score matching (Hyvärinen 2005)

$$\min_f \mathbb{E}_{p(x)} [\|\partial_x \log p(x) - \partial_x \log q(x)\|_2^2]$$

- ▶ Regularization reduces overfitting and also avoids spurious modes in  $q(x)$

We find  $f(\cdot)$  in a reproducing kernel Hubert space (RKHS)

- ▶ Kernel Lite (Strathmann et al. 2015)

$$f(x) = \sum_m \alpha_m k(x_m, x)$$

- ▶ **Deep kernel lite**

$$f(x) = \sum_m \alpha_m k_\theta(x_m, x)$$

$$k_\theta(x_m, x) = k(g_\theta(x_m), g_\theta(x))$$

## Goodness of fit test

- ▶ In benchmarking implicit models, it is difficult to evaluate performance using log likelihoods
- ▶ We would like to utilize an attractive property of the Stein operator  $\mathcal{T}_p$  acting on a function  $f$ . Define Stein operator as

$$\mathcal{T}_p[f] = \partial_x \log p(x) \cdot f(x) + \partial_x f(x)$$

then given another distribution  $q$  the following holds if and only if  $q = p$

$$\mathbb{E}_q [\mathcal{T}_p[f]] = 0$$

- ▶ So given data from distribution  $q$ , e.g. from a generative model
  - ▶ Variational autoencoder (VAE)
  - ▶ Generative adversarial network (GAN)
- ▶ And if the data are from a differentiable process resulting in true density  $p$ , then we can try to learn an operator  $\mathcal{T}_p$  to perform a statistical test. (Chwialkowski et al., 2016)
- ▶ Our approach
  - ▶ Construct a true model density  $p$
  - ▶ Train generative models to obtain  $p$
  - ▶ Evaluate  $\mathbb{E}_q [\mathcal{T}_p[f]]$  where  $f$  can be learned to maximise statistical power